

Representation of High Dimensional Data

Ross Goroshin



July 9, 2012




NEW YORK UNIVERSITY

Statistical Dependence and Natural Data Manifolds

From previous colloquium...




Dimensionality of Data & Dependence



- This illustration is representative of many processes
- However, dependence can be introduced without increasing the dimensionality
- Latent representation is NOT unique for generative processes of interest

Auto-Encoder Framework


- An Auto-Encoder(AE) is composed of an "encoder" and "decoder" (matrices in the simplest case)
- The encoder and decoder may correspond to completely different procedures (bases)
- The encoder transforms the data to latent space, and the decoder reconstructs the data from the latent representation
- AEs unify many data representation concepts and algorithms



Searching for structure. (Left) Three-dimensional data that are inherently one-dimensional. (Middle) A simple "autoencoder" network that is designed to compress three dimensions to one, through the narrow hidden layer of one unit. The inputs are labeled x, y, z , with outputs x', y' , and z' . (Right) A more complex autoencoder network that can represent highly nonlinear mappings from three dimensions to one, and from one dimension back out to three dimensions.

Example: $x = \cos(t), y = \sin(t), z = t$


Neural Network



hidden layer

output layer

Auto-Encoder Network



hidden layer

output layer

output layer

hidden layer

Unsupervised and Deep Learning

- Auto-encoders have been instrumental in the development of "Deep Learning"
- "Deep Learning" based techniques leverage the inherently hierarchical nature of our world
- Reconstruction cost is often used for unlabeled data
- Overcomes the so called "vanishing gradient" problem in neural networks
- Allows for 'wholesale' propagation of information, i.e. labels provide only a few bits for information
- Connections with maximizing likelihood, except that the partition function is missing
- $p(x)$ usually tells you almost everything you want to know, e.g. $p(y|x)$ where y is some label

Simple Auto-Encoders


Let X be the input (image), W_e is the encoding weight matrix (basis), f is an arbitrary nonlinear function applied point-wise, and X_r is called the reconstruction.

$$X_r = W_d f(W_e X + b_e) + b_d$$
$$\min_{W_e, W_d, b_e, b_d} \sum_{X \in D} \frac{1}{2} \|X - X_r\|^2$$

- Interesting/useful representations are obtained when Auto-Encoders are forced learn to compress the data
- Informational bottleneck is established via W_e or f , or both
- Cost is minimized usually via *stochastic gradient descent*
- If f is identity and we enforce that $W_e = W_d^T$ then a nearly orthogonal basis is learned
- If $W_e \neq W_d^T$ then $W_d \approx W_e f^{-1}$ (...why?)


Linear AE and PCA

- Encoding and decoding simply corresponds to multiplication by W_e and W_d , respectively
- If W_e and W_d are full rank, then the AE has the capacity to learn the identity function
- Using s.g.d the largest eigenvectors of the covariance (or their mixtures) are learned first
- More interesting representations can be hoped for if some nonlinear function is applied to the encoder outputs
- AEs with nonlinearity can be stacked



Visualizing the Role of f


Assume that the data is normalized, i.e. lies on a sphere



$$\frac{\partial E}{\partial W_d} = (X - X_r) f(W_e X)^T$$
$$\frac{\partial E}{\partial W_e} = W_d (X - X_r) f'(W_e X)^T X$$

Information Bottleneck via Regularization: Contractive AE


- Other techniques of obtaining "interesting" representations is to regularize the latent representation
- Contrary to more traditional regularization directly on weights, these correspond to "generic prior hypotheses"
- $L_{CAE} = \sum_{x \in D_n} \|x - x_r\|_2^2 + \lambda \sum_{ij} \left(\frac{\partial h_j(x)}{\partial x_i} \right)^2$



- $\left(\frac{\partial h_j(x)}{\partial x_i}\right)^2 = (h'_j(x))^2 W_{e_{ij}}^2$
- Important to tie or normalize weights
- If there were no reconstruction objective then the penalty on the Jacobian would produce a constant representation for all inputs either by saturating or weight decay
- However nearby images on the manifold must be reconstructed as distinct images
- The contractive pressure is counteracted by the reconstruction gradient in the directions tangent to the manifold
- Contractive penalty has a component in the direction of curvature gradient
- Can be interpreted roughly as curvature regularization

Visualizing the Effect of (Good) Regularization

- The data manifold (distribution) is depicted as yellow points
- Brightness is proportional to reconstruction error (i.e. darker areas are better reconstructed)
- Implicitly parameterizes the data manifold
- Regularization is analogous to the partition function in max likelihood models




No Regularization

With Regularization

SATAE: Regularization via Saturation


- Another regularization strategy is to encourage "activations" (h_i) in the saturation regime of the non-linearity (a.k.a activation function)
- In all cases, this should bound the volume of space which is well reconstructed



$$\text{Let } S = \{z \mid f'(z) = 0\}$$
$$f_c(z) = \inf_{z' \in S} |z - z'|$$


SATAE with *shrink()* Nonlinearity

- The saturation penalty exactly corresponds to an L_1 penalty
- This auto-encoder strongly resembles sparse coding using PSD with *shrink()* nonlinearity
- Important to enforce $\|w_d\|_2 = 1$



Experiments on Images

airplane



automobile



bird



cat




deer



dog




frog



horse




ship




truck




SATAE-shrink



SATAE-saturated linear




SATAE-saturated linear Progressively Increasing Saturation





Denoising AE

- This AE is trained to reconstruct an image from a corrupted version of itself
- The corruption process randomly chooses some proportion of the pixels in the image and sets them to zero
- This explicitly forces the AE to learn and exploit dependencies between pixels in images
- Can be interpreted as simultaneously learning $p(x|\tilde{x})$ and how to infer x from \tilde{x}
- If $\dim(Y) < \dim(X)$ then " $Y = f(X)$ is a representation of X which is well suited for capturing the main variations in the data, i.e. on the manifold"




Denoising AE



(a) No destroyed inputs

(b) 25% destruction

(c) 50% destruction



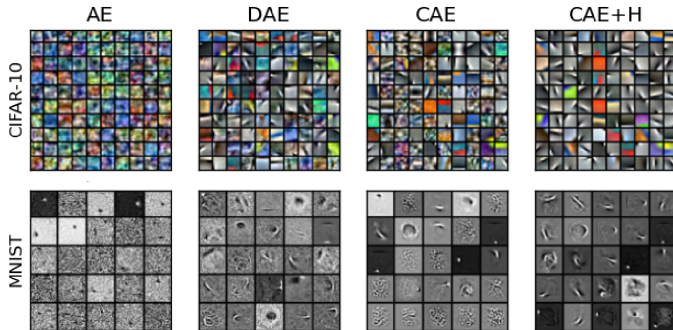
(d) Neuron A (0%, 10%, 20%, 50% destruction)



(e) Neuron B (0%, 10%, 20%, 50% destruction)

Relationship with other Auto-Encoders


- Denoising auto-encoders also achieve a contractive mapping of the *reconstruction* via a stochastic process, which indirectly encourages invariance in the latent representation
- Sparse auto-encoders encourage the activations to be zero which occurs in the saturated part of the nonlinearity (assuming a proper choice of nonlinearity, e.g. LISTA)



Enforcing a Factorization: Transforming Auto-Encoders

- "Move to a space in which the common variabilities can be described by linear transformations" (Hinton, 2012)
- In the case of objects in images: A representation in which the identity of the object (what?) is separately represented from the configuration (where?) would satisfy the above requirement
- This corresponds to a physically interpretable latent representation


Enforcing a Factorization: Transforming Auto-Encoders



Each capsule has the capacity to represent only one instance of its visual entity at a time

Inverting Auto-Encoder

- $R_W(G(y'_i)) = y''_i$ if $R_W = G^{-1}$ then $y'_i = y''_i$ thus we want to $\min_W \|y'_i - y''_i\|$



Algorithm for training a recognition network R_w parameterized by weight vector w :

Given: Training set X of n data vectors $\{x_1, x_2, \dots, x_n\}$, generative black box G , prototype code vector p .

Initialization: Set output biases of R_w using p , and the remaining weights to samples from a zero-mean Gaussian with a small standard deviation.




Weight update computed using the i^{th} (unlabelled) training case x_i :

Let y_i be the code vector inferred from x_i using the current recognition network R_w .

1. $y_i = R_w(x_i)$.
2. Perturb y_i randomly to create y'_i . (Note: The exact perturbation method is specified later.)
3. $x'_i = G(y'_i)$.
4. Supervised learning on (x'_i, y'_i) :
 - (a) $y''_i = R_w(x'_i)$.
 - (b) $E = \|y'_i - y''_i\|^2$.
 - (c) $w \leftarrow w - \eta \frac{\partial E}{\partial w}$.

Fig. 1. Summary of the breeder learning algorithm.

Inverting Auto-Encoder



Thank You

THE END